Dali Abid

Swedish citizen | Eligible to work in the EU | daliabid.se@gmail.com
linkedin.com/in/dali-abid | github.com/adeilm | adeilm.github.io

Summary

I am a Data Engineering student with a strong drive to specialize in Artificial Intelligence, seeking a 4–6-month internship starting in mid-January 2027. Through three internships over the past year, I have gained hands-on experience in Generative AI, particularly in RAG architectures (graph-based, hybrid, and multi-modal retrieval), LLM fine-tuning research, and the development of AI agents, agent harnesses, and orchestration workflows.

Education

Faculty of Sciences of Sfax — Tunisia

- Engineering Degree in Data Engineering — September 2024 – expected June 2027
- Integrated Pre-Engineering Studies — September 2022 – June 2024

Work and Research Experience

Themison – clinical trial management platform — August 2026 – November 2026
AI Engineer (part-time, alongside studies) — Copenhagen, Denmark (remote)

- Built a multimodal RAG-based Q&A system for long PDFs using hybrid retrieval with PostgreSQL/pgvector, providing precise answers with visual references highlighting the exact locations in the source documents.
- Built a document ingestion pipeline using Docling to handle different document layouts, preserve complex tables, enable row-level retrieval linked to parent tables, and provide citations pointing to the relevant pages and sections.
- Designed the LLM orchestration layer of a RAG assistant: query routing with a typed classifier, follow-up rewriting, and confidence thresholds with safe fallbacks, retries and timeouts; 94% routing accuracy on a hand-labelled test set.
- Built a bounded tool-calling agent that gathers linked evidence before the answer is written.
- Built an AI SQL agent that enables natural-language analysis of clinical trial data in multi-tenant PostgreSQL, using a semantic layer to understand study entities, relationships, and metrics while enforcing query validation and tenant isolation.

Zetabox — July 2026 – August 2026
AI Research and LLM Fine-Tuning Intern — Remote

- Evaluated LLM adaptation techniques (SFT, LoRA/QLoRA, RLHF, DPO, GRPO, distillation, and model merging), creating a practical decision guide to help engineering teams select approaches based on data requirements, computational resources, and project needs.
- Built the first version of a citation-backed RAG assistant that recommends LLM adaptation strategies based on client objectives, available data, infrastructure, and budget, making the research reusable for future projects.

BIMLO — July 2026
AI and Backend Engineering Intern — Remote

- Built a multimodal RAG system for engineering documents, combining PostgreSQL-based metadata filtering with scoped semantic search in ChromaDB to retrieve relevant information from natural-language questions.
- Developed a repeatable ingestion pipeline for engineering calculation documents (Word and PDF), organizing content by project, site, and revision to enable context-aware retrieval and document traceability.

Farm Pro Academy – e-learning platform — February 2026 – April 2026
Freelance Full-Stack Developer — Remote

- Developed a trilingual e-learning platform for agricultural training, with dedicated student, instructor, and admin portals.
- Built the full-stack application using Next.js, React, Node.js/Express, and MySQL, implementing role-based access and automated email notifications through Brevo.
- Deployed the platform to production on an OVH Linux server, configuring NGINX, PM2, and HTTPS with Certbot.

DK-Soft — July 2025 – August 2025
Full-Stack and Computer Vision Intern — Sfax, Tunisia

- Developed a virtual glasses try-on application using MediaPipe for real-time face tracking and Three.js/WebGL for 3D rendering, supporting both live-camera and photo modes.
- Built a Spring Boot backend integrating Spring AI with a MySQL product catalog for LLM-powered eyewear recommendations, alongside OCR-based prescription extraction.

Clinisys — June 2025
Data Analytics Intern — Sfax, Tunisia

- Developed a Power BI analytics dashboard backed by SQL Server and a lightweight data warehouse to analyze profitability and product flows, with automated KPI reporting.

Projects

AI-assisted intrusion detection

- Built and deployed a SOC enrichment prototype across five Ubuntu VMs in the faculty datacenter using CloudStack/KVM, Wazuh, Suricata and Elastic Stack.
- Developed a FastAPI service integrating Qwen3 14B via Ollama to explain alerts, suggest MITRE ATT&CK mappings, and generate Kibana investigation queries and remediation recommendations with JSON schema validation.
- Automated five end-to-end attack scenarios; measured a 63% reduction in mean alert comprehension time (342 to 127 seconds) in a controlled evaluation of 15 alerts with two student analysts.

AI-Assisted Tender Intelligence and Proposal Generation

- Built a tender intelligence platform combining n8n agents for tender discovery and buyer research with a FastAPI hybrid-retrieval RAG service matching requirements against CVs and company records.
- Implemented six LLM agent roles to generate proposal content, architecture diagrams, delivery plans and catalog-based budgets, with bounded review loops, slide regeneration and a browser editor supporting editable PowerPoint and PDF export.

Self-hosted Kubernetes and CI/CD platform

- Provisioned a three-node Kubernetes 1.29 lab using Vagrant and Ansible, configuring containerd, Calico networking and NFS-backed persistent storage.
- Built a Jenkins pipeline to build Docker images, publish them to Nexus and deploy a three-tier application with an NGINX frontend, Express.js API and PostgreSQL database to Kubernetes.
- Configured application service accounts and RBAC, network policies and baseline Pod Security Admission, alongside health probes and resource limits for frontend and backend workloads.

Skills

Generative AI & agents: Hybrid and multimodal RAG, GraphRAG, LangChain, LangGraph, n8n, tool-calling agents, SQL agents, prompt engineering, Ollama

Retrieval & evaluation: BM25, multilingual embeddings, cross-encoder reranking, retrieval benchmarking, LLM-as-a-judge evaluation, PostgreSQL/pgvector, ChromaDB, Neo4j

Programming & backend: Python, TypeScript, SQL, FastAPI, Node.js/Express, Spring Boot, REST APIs

Machine learning & document processing: PyTorch, TensorFlow, scikit-learn, Pandas, NumPy, Hugging Face, Sentence-Transformers, Docling, OCR

LLM adaptation (research): SFT, LoRA/QLoRA, RLHF, DPO, GRPO, distillation, model merging

Cloud & DevOps: Linux, Docker, Kubernetes, Azure, Ansible, Vagrant, Jenkins, Nexus, NGINX, Git

Languages: English and French: fluent; Arabic: native; Swedish: beginner

Awards and Activities

1st place: Master Faster Hackathon – GDSC FSS, 2025.

1st place: REGEN-X – IEEE FSS, 2024.

2nd place: CSTAM Technical Challenge – IEEE ENICarthage, 2025; Tech4Health – IEEE ENET'COM, 2025.

Leadership: Co-founder and technical lead, CPC competitive programming club (2023–2025).

Memberships: IEEE and GDSC.